

Problem Statement: -

Solar power consumption has been recorded by city councils at regular intervals. The reason behind doing so is to understand how businesses are using solar power so that they can cut down on nonrenewable sources of energy and shift towards renewable energy. Based on the data, build a forecasting model and provide insights on it.

In [1]:

```
#importing required libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

In [2]:

```
solar = pd.read_csv("D:/DataScience/Class/assignment working/Forecasting/solarpower_cumuldaybyday2.csv")
```

In [3]:

```
solar["date"] = pd.to_datetime(solar["date"],infer_datetime_format = True)
indexed_data = solar.set_index(["date"])
```

In [4]:

```
from datetime import datetime
indexed_data.head(10)
```

Out[4]:

cum_power	
date	
2011-10-26	0.1
2011-10-27	10.2
2011-10-28	20.2
2011-10-29	29.6
2011-10-30	34.2
2011-10-31	38.0
2011-11-01	46.6
2011-11-02	51.6
2011-11-03	58.6
2011-11-04	60.5

In [5]:

```
indexed_data.tail()
```


Out[5]:

cum_power	
date	
2018-10-22	28101.0
2018-10-23	28109.0
2018-10-24	28115.0
2018-10-25	28117.0
2018-10-26	28120.0

as the data is from 2011-10-26 to 2018 -10 -26 the cycle of 360 days startrts at 26th and ends at next 36th so no need to drop any rows

In [6]:

```
#plotting Graph
plt.xlabel("Date")
plt.ylabel("Power Consumption")
plt.plot(indexed_data)
plt.show()
```



Cheking Stationarity of data

Checking Stationarity of data

ths is the default requirement for any forecasting algorithm ,the data should be stationary

```
In [7]:
pd.options.display.max_rows = 999

In [8]:
pd.options.display.max_columns = 100

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [10]:
#checking stationarity with dicky-fuller test
from statsmodels.tsa.stattools import adfuller

dfctest = adfuller(indexed_data["cum_power"], autolag = "AIC")
adfoutput = pd.Series(dfctest[0:4],index = ["test statastic", "p_value","lag_used","number of observations used"])
for key,value in dfctest[4].items():
    adfoutput["Critical value (%s)"%key]=value

print(adfoutput)

test statastic          -0.421484
p_value                 0.906449
lag_used                20.000000
number of observations used  2537.000000
Critical value (1%)      -3.432930
Critical value (5%)      -2.862680
Critical value (10%)     -2.567377
dtype: float64
```

we can not reject null hypothesis so data is not stationary

null hypothesis = data is not stationary

```
In [ ]:


In [12]:
solar.isna().sum()


Out[12]:
date          0
cum_power     0
dtype: int64

In [14]:
from statsmodels.graphics.tsaplots import plot_acf ,plot_pacf

In [15]:
#checking accuracy with different lags and differentiation and moving average

In [18]:
#normal acf and pacf plot
plot_acf(indexed_data)
plot_pacf(indexed_data)
plt.show()
```





In [20]:


```
# first differencing
indexed_data["first_diff"] = indexed_data["cum_power"] - indexed_data["cum_power"].shift(1)
```

In [33]:

```
plt.plot(indexed_data.first_diff)
```

Out[33]:



[<matplotlib.lines.Line2D at 0x208dcd0d730>]



In [34]:

```
plot_acf(indexed_data.first_diff)
plot_pacf(indexed_data.first_diff)
plt.show()
```

C:\Users\theas\anaconda3\lib\site-packages\numpy\core_asarray.py:83: UserWarning: Warning: converting a masked element to nan.
return array(a, dtype, copy=False, order=order)



In [36]:

```
adfuller(indexed_data.first_diff.dropna())
```

Out[36]:

```
(-3.0856724129568414,
0.02763182430737851,
19,
2537,
{'1%': -3.4329301847920486,
'5%': -2.862679919243664,
'10%': -2.5673768219208686},
15175.367039787136)
```

In [41]:


```
# second differencing
indexed_data["second_diff"] = indexed_data["cum_power"] - indexed_data["cum_power"].shift(1).shift(1)
```

In [42]:

```
plt.plot(indexed_data["second_diff"])
```

Out[42]:

[<matplotlib.lines.Line2D at 0x208dc8b00d0>]



In [43]:

```
indexed_data["second_diff"].head()
```

Out[43]:

```
date
2011-10-26    NaN
2011-10-27    NaN
2011-10-28    20.1
2011-10-29    19.4
2011-10-30    14.0
Name: second_diff, dtype: float64
```

In [44]:

```
adfuller(indexed_data.second_diff.dropna())
```

Out[44]:

```
(-2.965202630211606,
 0.0382743102354102,
 26,
 2529,
 {'1%': -3.432938355012086,
  '5%': -2.8626835272597217,
  '10%': -2.567378742868999},
 15251.943584943776)
```

In [45]:

```
indexed_data["second_diff"]
```

Out[45]:

```
date
2011-10-26    NaN
2011-10-27    NaN
2011-10-28    20.1
2011-10-29    19.4
2011-10-30    14.0
...
2018-10-22    15.0
2018-10-23    14.0
2018-10-24    14.0
2018-10-25     8.0
2018-10-26     5.0
Name: second_diff, Length: 2558, dtype: float64
```

In []:

In [53]:

```
#pip install pmdarima
```

In [55]:

```
from pmdarima import auto_arima
```

In [56]:

```
auto_model=auto_arima(indexed_data.cum_power,trace=True) ### returns best p,d,q values based on AIC score
```

```
Performing stepwise search to minimize aic
ARIMA(2,1,2) (0,0,0) [0] intercept      : AIC=15320.006, Time=2.34 sec
ARIMA(0,1,0) (0,0,0) [0] intercept      : AIC=18025.637, Time=0.07 sec
ARIMA(1,1,0) (0,0,0) [0] intercept      : AIC=15947.568, Time=0.30 sec
ARIMA(0,1,1) (0,0,0) [0] intercept      : AIC=16895.797, Time=0.43 sec
ARIMA(0,1,0) (0,0,0) [0]                : AIC=20652.126, Time=0.05 sec
ARIMA(1,1,2) (0,0,0) [0] intercept      : AIC=15319.616, Time=1.15 sec
ARIMA(0,1,2) (0,0,0) [0] intercept      : AIC=16517.081, Time=0.69 sec
ARIMA(1,1,1) (0,0,0) [0] intercept      : AIC=15379.388, Time=0.88 sec
```

ARIMA(1,1,3) (0,0,0) [0] intercept : AIC=15320.898, Time=1.48 sec
ARIMA(0,1,3) (0,0,0) [0] intercept : AIC=16248.109, Time=1.18 sec
ARIMA(2,1,1) (0,0,0) [0] intercept : AIC=15318.019, Time=1.27 sec
ARIMA(2,1,0) (0,0,0) [0] intercept : AIC=15701.892, Time=0.43 sec
ARIMA(3,1,1) (0,0,0) [0] intercept : AIC=15320.013, Time=1.87 sec
ARIMA(3,1,0) (0,0,0) [0] intercept : AIC=15535.761, Time=0.67 sec
ARIMA(3,1,2) (0,0,0) [0] intercept : AIC=15320.532, Time=3.22 sec
ARIMA(2,1,1) (0,0,0) [0] : AIC=15323.620, Time=0.64 sec

Best model: ARIMA(2,1,1) (0,0,0) [0] intercept
Total fit time: 16.754 seconds

In [57]:

```
auto_model.summary()
```

Out[57]:

SARIMAX Results

Dep. Variable:	y	No. Observations:	2558			
Model:	SARIMAX(2, 1, 1)		Log Likelihood	-7654.009		
Date:	Sun, 04 Jul 2021		AIC	15318.019		
Time:	14:48:18		BIC	15347.252		
Sample:	0		HQIC	15328.620		
- 2558						
Covariance Type:	opg					
	coef	std err	z	P> z	[0.025	0.975]
intercept	0.0635	0.044	1.447	0.148	-0.023	0.149
ar.L1	1.1941	0.020	59.228	0.000	1.155	1.234
ar.L2	-0.2001	0.019	-10.373	0.000	-0.238	-0.162
ma.L1	-0.8640	0.012	-73.828	0.000	-0.887	-0.841
sigma2	23.2928	0.582	40.034	0.000	22.152	24.433
Ljung-Box (L1) (Q):	0.00	Jarque-Bera (JB):	60.60			
Prob(Q):	1.00	Prob(JB):	0.00			
Heteroskedasticity (H):	1.10	Skew:	0.01			
Prob(H) (two-sided):	0.17	Kurtosis:	3.75			

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [171]:

```
df=pd.read_csv("D:/DataScience/Class/assignment working/Forcasting/solarpower_cumuldaybyday2.csv")
```

Building ARIMA model

In [174]:

```
Train = df.head(2192)  
Test=df.tail(366)
```

In [175]:

```
# Creating a function to calculate the MAPE value for test data  
def MAPE(pred,org):  
    temp = np.abs((pred-org)/org)*100  
    return np.mean(temp)
```

In [184]:

```
# Simple Exponential Method  
from statsmodels.tsa.holtwinters import SimpleExpSmoothing  
ses_model = SimpleExpSmoothing(Train["cum_power"]).fit(smoothing_level=1)  
pred_ses = ses_model.predict(start = Test.index[0], end = Test.index[-1])  
MAPE(pred_ses, Test.cum_power)
```

Out[184]:

5.922982997009587

In [189]:

```
np.sqrt(np.mean(pred_ses-Test.cum_power)**2)
```

Out[189]:

1584.4754098360656

In [186]:

```
# Holt method  
from statsmodels.tsa.holtwinters import Holt  
hw_model = Holt(Train["cum_power"]).fit(smoothing_level=0.01)  
pred_hw = hw_model.predict(start = Test.index[0], end = Test.index[-1])  
MAPE(pred_hw, Test.cum_power)
```

C:\Users\theas\anaconda3\lib\site-packages\statsmodels\tsa\holtwinters\model.py:427: FutureWarning: After 0.13 initialization must

```
be handled at model creation
warnings.warn(
C:\Users\theas\anaconda3\lib\site-packages\statsmodels\tsa\holtwinters\model.py:920: ConvergenceWarning: Optimization failed to co
nverge. Check mle_retvals.
warnings.warn(
```

Out[186]:

2.327441434347866

In [188]:

```
np.sqrt(np.mean(pred_hw-Test.cum_power)**2)
```

Out[188]:

576.2850881654938

In [193]:

```
# Holts winter exponential smoothing with additive seasonality and additive trend
from statsmodels.tsa.holtwinters import ExponentialSmoothing
hwe_model_add_add = ExponentialSmoothing(Train["cum_power"], seasonal = "add", trend = "add", seasonal_periods = 2).fit(smoothing_level=0.01)
pred_hwe_add_add = hwe_model_add_add.predict(start = Test.index[0], end = Test.index[-1])
MAPE(pred_hwe_add_add, Test.cum_power)
```

Out[193]:

2.0681153138377666

In [195]:

```
np.sqrt(np.mean(pred_hwe_add_add-Test.cum_power)**2)
```

Out[195]:

427.2307160122235

In [198]:

```
# Holts winter exponential smoothing with multiplicative seasonality and additive trend
hwe_model_mul_add = ExponentialSmoothing(Train["cum_power"], seasonal = "mul", trend = "add", seasonal_periods = 2).fit(smoothing_level=0.01)
pred_hwe_mul_add = hwe_model_mul_add.predict(start = Test.index[0], end = Test.index[-1])
MAPE(pred_hwe_mul_add, Test.cum_power)
```

Out[198]:

1.9882712951755088

In [199]:

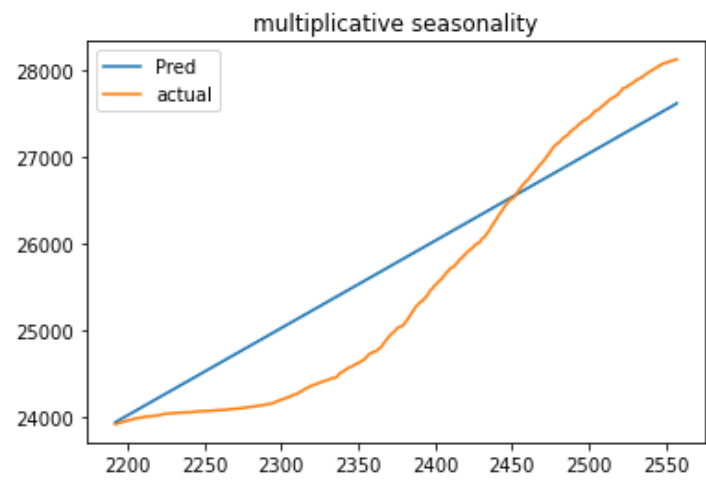
```
np.sqrt(np.mean(pred_hwe_mul_add-Test.cum_power)**2)
```

Out[199]:

275.19585296625985

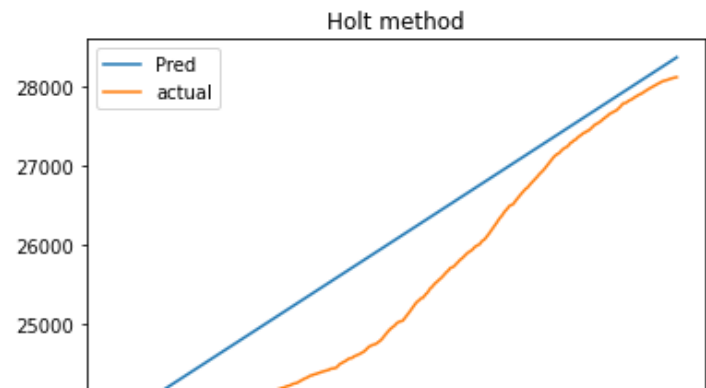
In [203]:


```
plt.plot(pred_hwe_mul_add,label="Pred")
plt.plot(Test.cum_power,label="actual")
plt.title("multiplicative seasonality")
plt.legend()
plt.show()
```



In [204]:

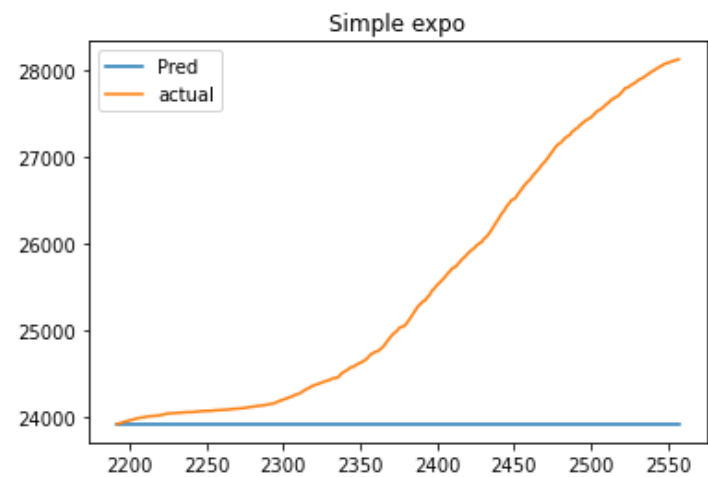
```
plt.plot(pred_hw,label="Pred")
plt.plot(Test.cum_power,label="actual")
plt.title("Holt method")
plt.legend()
plt.show()
```





In [205]:

```
plt.plot(pred_ses,label="Pred")
plt.plot(Test.cum_power,label="actual")
plt.title("Simple expo")
plt.legend()
plt.show()
```



Final Model

predicting for next one year

In [206]:

```
pred = pd.read_csv("D:/DataScience/Class/assignment working/Forcasting/pred_solarpower_cumuldaybyday2.csv")
```

In [221]:

```
final_mod = ExponentialSmoothing(solar["cum_power"], seasonal = "mul", trend = "add", seasonal_periods = 2).fit()
```

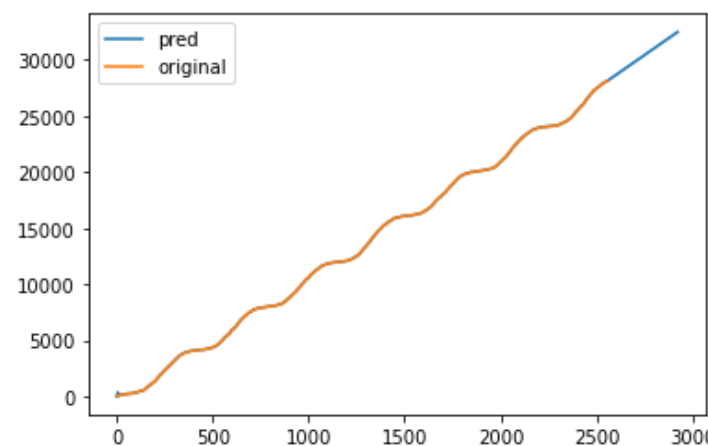
C:\Users\theas\anaconda3\lib\site-packages\statsmodels\tsa\holtwinters\model.py:920: ConvergenceWarning: Optimization failed to converge. Check mle_retvals.
warnings.warn(

In [222]:

```
final_pred =final_mod.predict(start = pred.index[0], end = pred.index[-1])
```

In [223]:

```
plt.plot(final_pred,label="pred")
plt.plot(solar.cum_power,label="original")
plt.legend()
plt.show()
```



In [224]:

```
import statsmodels.api as sm
```

In [225]:

```
model = sm.tsa.statespace.SARIMAX(solar.cum_power,order=(1,1,1),seasonal_order=(1,1,1,2)) #seasonal_order=(p,d,q,m) #m=lag
results=model.fit()
```


In [226]:

```
predic=results.predict(start = pred.index[0], end = pred.index[-1])
```

In [227]:

```
plt.plot(predic,label="pred")
plt.plot(solar.cum_power,label="original")
plt.legend()
plt.show()
```





Summary and inference

Im anablw to solve this problem ,as you can see the results are pathetic, i need model solution for this problem from your side,thank you

In []: